

What is Service Mesh and Istio

A **service mesh** is a dedicated infrastructure layer for making service-to-service communication safe, fast, and reliable.

Istio is a service mesh which allows you to connect, manage and secure your microservices in an easy and none intrusive way.

Some of the features that offer Istio are:

- Intelligent routing and load balancing
- Resilience against network failures
- Policy enforcement between services
- Observability of your architecture. Tracing and Metrics
- Securing service to service communication

Istio Architecture

Istio is composed of two major components:

- **Data plane** which is composed of **Envoy** proxies deployed as sidecar container along with your service for managing network along with policy and telemetry features.
- **Control plane** which is in charge of managing and configuring all **Envoy** proxies.

All communication within your **service mesh** happens through **Envoy** proxy, so any network logic to apply is moved from your service into your infrastructure.

Key Concepts of Istio

DestinationRule

A **DestinationRule** configures the set of rules to be applied when forwarding traffic to a service. Some of the purposes of a **DestinationRule** are describing circuit breakers, load balancer, and TLS settings or define **subsets** (named versions) of the destination host so they can be reused in other Istio elements.

For example to define two services based on the version label of a service with hostname **recommendation** you could do:

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: recommendation
  namespace: tutorial
spec:
  host: recommendation
  subsets:
  - labels:
    version: v1
    name: version-v1
  - labels:
    version: v2
    name: version-v2
```

VirtualService

A **VirtualService** describes the mapping between one or more user-addressable destinations to the actual destination inside the mesh.

For example, to define two virtual services where the traffic is split between 50% to each one.

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: recommendation
  namespace: tutorial
spec:
  hosts:
  - recommendation
  http:
  - route:
    - destination:
        host: recommendation
        subset: version-v1
      weight: 90
    - destination:
        host: recommendation
        subset: version-v2
      weight: 10
```

ServiceEntry

A **ServiceEntry** is used to configure traffic to external services of the mesh such as APIs or legacy systems. You can use it in conjunction with a **VirtualService** and/or **DestinationRule**.

For example to configure *httpbin* external service:

```
apiVersion: networking.istio.io/v1alpha3
kind: ServiceEntry
metadata:
  name: httpbin-egress-rule
  namespace: istioegress
spec:
  hosts:
  - httpbin.org
  ports:
  - name: http-80
    number: 80
    protocol: http
```

Gateway

A **Gateway** is used to describe a load balancer operating at the edge of the mesh for incoming/outgoing HTTP/TCP connections. You can bind a **Gateway** to a **VirtualService**.

To configures a load balancer to allow external https traffic for host foo.com into the mesh:

```
apiVersion: networking.istio.io/v1alpha3
kind: Gateway
metadata:
  name: foo-gateway
spec:
  servers:
  - port:
      number: 443
      name: https
      protocol: HTTPS
    hosts:
    - foo.com
    tls:
      mode: SIMPLE
      serverCertificate: /tmp/tls.crt
      privateKey: /tmp/tls.key
```

Getting started with Istio

Istio can be installed with *automatic sidecar injection* or without it. We recommend as starting point **without** *automatic sidecar injection* so you understand each of the steps.

Installing Istio

First you need to download Istio and register in PATH:

```
open https://github.com/istio/istio/releases/

cd istio-1.0.2
export ISTIO_HOME=`pwd`
export PATH=$ISTIO_HOME/bin:$PATH
```

You can install Istio into Kubernetes cluster by either using helm install or helm template.

```
$ helm template install/kubernetes/helm/istio \
  --name istio --namespace istio-system \
  --set sidecarInjectorWebhook.enabled=false \
  > $HOME/istio.yaml

kubectl create namespace istio-system
kubectl create -f $HOME/istio.yaml
```

Wait until all pods are up and running.

Intelligent Routing

Routing some percentage of traffic between two versions of recommendation service:

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: recommendation
  namespace: tutorial
spec:
  hosts:
  - recommendation
  http:
  - route:
    - destination:
        host: recommendation
        subset: version-v1
      weight: 75
    - destination:
        host: recommendation
        subset: version-v2
      weight: 25
```

Routing to a specific version in case of prefixed URI and cookie with a value matching a regular expression:

```
spec:
  hosts:
  - ratings
  http:
  - match:
    - headers:
        cookie:
          regex: "^(.*?;)?(user=jason)(;.*)?"
        uri:
          prefix: "/ratings/v2/"
    route:
    - destination:
        host: ratings
        subset: version-v2
```

Possible **match** options:

| Field | Type | Description |
|--------------|--------------------------|--|
| uri | StringMatch | URI value to match. exact, prefix, regex |
| scheme | StringMatch | URI Scheme to match. exact, prefix, regex |
| method | StringMatch | Http Method to match. exact, prefix, regex |
| authority | StringMatch | Http Authority value to match. exact, prefix, regex |
| headers | map<string, StringMatch> | Headers key/value. exact, prefix, regex |
| port | int | Set port being addressed. If only one port exposed, not required |
| sourceLabels | map<string, string> | Caller labels to match |
| gateways | string[] | Names of the gateways where rule is applied to. |

Sending traffic depending on caller labels:

```
- match:
  - sourceLabels:
      app: preference
      version: v2
    route:
  - destination:
      host: recommendation
      subset: version-v2
- route:
  - destination:
      host: recommendation
      subset: version-v1
```

When caller contains labels `app=preference` and `version=v2` traffic is routed to **subset** `version-v2` if not routed to `version-v1`

Mirroring traffic between two versions:

```
spec:
  hosts:
  - recommendation
  http:
  - route:
      - destination:
          host: recommendation
          subset: version-v1
    mirror:
      host: recommendation
      subset: version-v2
```

For routing purposes `VirtualService` also supports **redirects**, **rewrites**, **corsPolicies** or **appending** custom headers.

Apart from HTTP rules, `VirtualService` also supports matchers at *tcp* level.

```
spec:
  hosts:
  - postgresql
  tcp:
  - match:
      - port: 5432
        sourceSubnet: "172.17.0.0/16"
    route:
      - destination:
          host: postgresql
          port:
            number: 5555
```

Possible **match** options at *tcp* level:

| Field | Type | Description |
|--------------------------|---------------------|--|
| destinationSubnet | string | IPv4 or IPv6 of destination with optional subnet |
| port | int | Set port being addressed. If only one port exposed, not required |
| sourceSubnet | string | IPv4 or IPv6 of source with optional subnet |
| sourceLabels | map<string, string> | Caller labels to match |
| gateways | string[] | Names of the gateways where rule is applied to |

Resilience

Retry 3 times when things go wrong before throwing the error upstream.

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: recommendation
  namespace: tutorial
spec:
  hosts:
  - recommendation
  http:
  - retries:
      attempts: 3
      perTryTimeout: 4.000s
    route:
      - destination:
          host: recommendation
          subset: version-v1
```

You can add timeouts to communications, for example aborting call after 1 second:

```
http:
- route:
  - destination:
      host: recommendation
      timeout: 1.000s
```

If the request is forwarded to a certain instance and it fails (e.g. returns a 50x error code), then this instance of an instance/pod is ejected to serve any other client request for an amount of time. In next example there must occur 5 consecutive errors before pod is ejected, ejection analysis occurs every 15 seconds, in case of ejection host will be ejected for 2 minutes and any host can be ejected.

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: recommendation
  namespace: tutorial
spec:
  host: recommendation
  trafficPolicy:
    outlierDetection:
      baseEjectionTime: 2m
      consecutiveErrors: 5
      interval: 15.000s
      maxEjectionPercent: 100
  subsets:
```

`trafficPolicy` can be applied at subset level to make it specific to a subset instead of all them.

You can also create connection pools at *tcp* and *http* level:

```
trafficPolicy:
  connectionPool:
    http:
      http1MaxPendingRequests: 100
      http2MaxRequests: 100
      maxRequestsPerConnection: 1
    tcp:
      maxConnections: 100
      connectTimeout: 50ms
```

Traffic Policy possible values:

| Field | Type | Description |
|-------------------------|------------------------|---|
| loadbalancer | LoadBalancerSettings | Controlling load blancer algorithm |
| connectionPool | ConnectionPoolSettings | Controlling connection pool |
| outlierDetection | OutlierDetection | Controlling eviction of unhealthy hosts |

| Field | Type | Description |
|--------------------------|---------------------|---|
| | | |
| tls | TLSSettings | TLS settings for connections |
| | | |
| portLevelSettings | PortTrafficPolicy[] | Traffic policies specific to concrete ports |

Policy Enforcement

Istio provides a model to enforce authorization policies in the communication between policies. You can, for example, black-list or white-list intercommunication between services or add some quota.

You can configure that preference service only allows requests from the recommendation service.

```
apiVersion: "config.istio.io/v1alpha2"
kind: listchecker
metadata:
  name: preferencewhitelist
spec:
  overrides: ["recommendation"]
  blacklist: false
```

```
apiVersion: "config.istio.io/v1alpha2"
kind: listentry
metadata:
  name: preferencesource
spec:
  value: source.labels["app"]
---
apiVersion: "config.istio.io/v1alpha2"
kind: rule
metadata:
  name: checkfromcustomer
spec:
  match: destination.labels["app"] == "preference"
  actions:
  - handler: preferencewhitelist.listchecker
    instances:
    - preferencesource.listentry
```

Source part is configured by using `listchecker` (to provide the list of allows hosts) and `listentry` (to configure how to get whitelist value from the request) elements. Destination part and rule is configured by using the `rule` element.

| Field | Type | Description |
|--------------------|--------|---|
| providerUrl | string | Url where to load the list to check against, can be empty |

| Field | Type | Description |
|-------------------------|---------------|--|
| refreshInterval | Duration | How often provider is polled |
| | | |
| ttl | Duration | How long keep list before discarding it |
| | | |
| catchingInterval | Duration | How long a caller can cache an answer befoer ask again |
| | | |
| catchingUseCount | int | Number of times a caller can use a cached answer |
| | | |
| overrides | string[] | List of entries consulted first before <code>providerUrl</code> |
| | | |
| entryType | ListEntryType | The kind (STRINGS, CASE_INSENSITIVE_STRINGS, IP_ADDRESSES, `REGEX) of list entry and overrides |
| | | |
| blacklist | boolean | the list operates as a blacklist or a whitelist |

Telemetry, Monitoring and Tracing

Isito comes with observability in mind providing out-of-the-box integration with Prometheus/Graphana and Jaeger.

Service to Service Security

You can secure the communication between all services by enabling mutual TLS (peer authentication).

First, you need to enable mutual TLS.

You can enable it globally:

```
apiVersion: "authentication.istio.io/v1alpha1"
kind: "MeshPolicy"
metadata:
  name: "default"
spec:
  peers:
  - mtls: {}
```

Or by namespace:

```
apiVersion: "authentication.istio.io/v1alpha1"
kind: "Policy"
metadata:
  name: "default"
  namespace: "tutorial"
spec:
  peers:
  - mtls: {}
```

Applying mTLS to specific destination and port:

```
spec:
  target:
  - name: preference
    ports:
    - number: 9000
```

If `ports` not set then it is applied to all ports.

| Field | Type | Description |
|-------------------------|----------------------------------|---|
| | | |
| peers | PeerAuthentication Method[] | List of authentication methods for peer auth |
| | | |
| peerIsOptional | boolean | Accept request when none of the peer authentication methods defined are satisfied |
| | | |
| targets | TargetSelector[] | Destinations where policy should be applied on. Enabled all by default |
| | | |
| origins | OriginAuthentication Method[] | List of authentication methods for origin auth |
| | | |
| originIsOptional | boolean | Accept request when none of the origin authentication methods defined are satisfied |

| Field | Type | Description |
|-------------------------|------------------|--|
| principalBinding | PrincipalBinding | Peer or origin identity should be use for principal. USE_PEER by default |

End user authentication (origin authentication) using JWT:

```
spec:
  origins:
  - jwt:
      issuer: "https://keycloak/auth/realms/istio"
      audiences:
      - "customer-tutorial"
      jwksUri: >
        https://keycloak/auth/realms/istio
          /protocol/openid-connect/certs
      principaBinding: USE_ORIGIN
```

At this time, Origins only support JWT. Possible values for JWT are:

| Field | Type | Description |
|-------------------|----------|--|
| issuer | string | Issuer of the token |
| audiences | string[] | List of JWT <i>audiences</i> allowed to access |
| jwksUri | string | URL of the public key to validate signature |
| jwtParams | string[] | JWT is sent in a query parameter |
| jwtHeaders | string[] | JWT is sent in a request header. If empty Authorization: Bearer \$token |

After enabling mTLS, you need to configure it at the client side by using a DestinationRule. Need to set which hosts communicate through mTLS using host field.

```
apiVersion: "networking.istio.io/v1alpha3"
kind: "DestinationRule"
metadata:
  name: "default"
  namespace: "tutorial"
spec:
  host: "*.tutorial.svc.cluster.local"
  trafficPolicy:
    tls:
      mode: ISTIO_MUTUAL
```

If **ISTIO_MUTUAL** is set, Istio configures client certificate, private key and CA crtificates with its internal implementation.

| Field | Type | Description |
|--------------------------|----------|--|
| httpsRedirect | boolean | Send 301 redirect when communication is using HTTP asking to use HTTPS |
| mode | TLSmode | How TLS is enforced. Values <i>PASSTHROUGH, SIMPLE, MUTUAL</i> |
| serverCertificate | string | The location to the file of the server-side TLS certificate |
| privateKey | string | The location to the file of the server's private key |
| caCertificates | string | The location to the file of the certificate authority certificates |
| subjectAltNames | string[] | Alternate names to verify the subject identity |

Istio RBAC

Istio’s authorization feature provides access control for services in an Istio Mesh.

To enable RBAC:

```
apiVersion: "rbac.istio.io/v1alpha1"
kind: RbacConfig
metadata:
  name: default
spec:
  mode: 'ON_WITH_INCLUSION'
  inclusion:
    namespaces: ["tutorial"]
```

Valid modes are: ON, OFF, ON_WITH_INCLUSION, ON_WITH_EXCLUSION. inclusion is used when WITH_INCLUSION and exclusion used when WITH_EXCLUSION. They support the next properties:

| Field | Type | Description |
|-------------------|----------|----------------------|
| services | string[] | A list of services |
| namespaces | string[] | A list of namespaces |

Granting access (**what**) to all services, when using the GET method and given destination services.

```
apiVersion: "rbac.istio.io/v1alpha1"
kind: ServiceRole
metadata:
  name: service-viewer
  namespace: tutorial
spec:
  rules:
  - services: ["*"]
    methods: ["GET"]
    constraints:
    - key: "destination.labels[app]"
      values: ["customer", "recommendation", "preference"]
```

| Field | Type | Description |
|--------------------|--------------|---------------------------------|
| services | string[] | List of service names to apply. |
| paths | string[] | List of HTTP paths |
| methods | string[] | List of HTTP methods |
| constraints | Constraint[] | Extra constraints |

And the `Constraint` is an array of pairs `key (string)` and `values (string[])`. Valid keys are:

| Key Example | Value Example |
|--|--|
| <code>destination.ip</code> | <code>["10.1.2.3", "10.2.0.0/16"]</code> |
| <code>destination.port</code> | <code>["80", "443"]</code> |
| <code>destination.labels[version]</code> | <code>["v1", "v2"]</code> |
| <code>destination.name</code> | <code>["productpage*"]</code> |
| <code>destination.namespace</code> | <code>["tutorial"]</code> |
| <code>destination.user</code> | <code>["customer-tutorial"]</code> |
| <code>request.headers[X-Custom-Token]</code> | <code>["345CFA3"]</code> |

Granting to all subjects (**who**) previous defined roles (**what**).

```
apiVersion: "rbac.istio.io/v1alpha1"
kind: ServiceRoleBinding
metadata:
  name: bind-service-viewer
  namespace: tutorial
spec:
  subjects:
  - user: "*"
  roleRef:
    kind: ServiceRole
    name: "service-viewer"
```


| Field | Type | Description |
|-------------------------|--------|------------------------------------|
| <code>user</code> | string | username/ID (Service Account). |
| <code>properties</code> | map | Properties to identify the subject |

Next properties are supported:


| Key Example | Value Example |
|--|---|
| <code>source.ip</code> | <code>"10.1.2.3"</code> |
| <code>source.namespace</code> | <code>"default"</code> |
| <code>source.principal</code> | <code>"customer"</code> |
| <code>request.headers[User-Agent]</code> | <code>"Mozilla/*"</code> |
| <code>request.auth.principal</code> | <code>"users.tutrial.org/654654"</code> |
| <code>request.auth.audiences</code> | <code>"tutorial.org"</code> |
| <code>request.auth.presenter</code> | <code>"654654.tutorial.org"</code> |
| <code>request.auth.claims[iss]</code> | <code>"*@redhat.com"</code> |

Last property refers to JWT claim named `iss`. Obviously, you can use any other claim for this purpose. Usually, you might use `group` claim to allow access to users under a specific group.

Authors :



@alexsotob
Java Champion and SW Engineer at Red Hat



@christianposta
Chief Architect at Red Hat

v1.0.0